

ISOM 3390: Business Programming in R

Instructor: Bingjie Qian (PhD candidate)

Summer 2020, HKUST

Topic 3: Control Structures

3.1 Overview of Control Structures

R programs are made up of **statements**. Statements are the basic units of instructions that the R interpreter parses and processes.

R statements are separated with a ; or a new line.

A semicolon always indicates the end of a statement.

```
x = 1:100; mean(x)
## [1] 50.5
```

A new line sometimes does not indicate the end of a statement. Statements are processed when they are syntactically complete.

```
course <-          # this line of code is syntactically incomplete
  c('R programming', 'Business analytics')
course

## [1] "R programming"      "Business analytics"

# try the code in the console
```

Statements can be grouped together using braces { and } to form a larger unit of execution, called a **block**. A block is processed as a whole. R only prints the output of the last statement.

```
{
x <- 1:100
mean(x)
course <- c('R programming', 'Business analytics')
course
}

## [1] "R programming"      "Business analytics"
```

In general, R executes statements sequentially from the first to the last statement, but sometimes we may want to control the flow of a program.

Control Structures

Control structures allow us to control the flow of a program. Controls structures in R are very similar to those of other programming languages.

Two common control structures in R:

- **Conditional Statements (Conditionals):** execute different sets of statements, depending on whether a condition is true or false.
- **Iterative Statements (Loops):** execute a set of statements repeatedly.

Most control structures are not used in interactive sessions (in the console), rather they are usually used for writing functions or longer expressions.

3.2 Conditional Statements

Conditional statements allow us to execute different sets of statements, depending on whether a condition is true or false.

3.2.1 if Statement




Fig. 1 Flowchart of if Statement

Syntax:

```
if (condition){  
  statements  
}
```

- If the condition is TRUE, the statements get executed.
- If the condition is FALSE, nothing happens.

```
x <- 6  
if (x > 0) {
```

```
print("A positive number")
}
## [1] "A positive number"
```

3.2.2 if else Statement




Fig. 2 Flowchart of if else Statement

Syntax:

```
if (condition) {
  statements_1
} else {
  statements_2
}
```

- If the condition is TRUE, statements_1 get executed.
- If the condition is FALSE, statements_2 get executed.

```
x <- 6
if (x > 100) {
  print("A large number")
} else {
  print("A small number")
}
## [1] "A small number"
```

3.2.3 if else Ladder

The if else ladder allows us to choose among more than 2 alternatives. Only one of the alternatives will be executed.




Fig. 3 Flowchart of if else Ladder

Syntax:

```

if (condition_1) {
    statements_1
} else if (condition_2) {
    statements_2
} else {
    statements_3
}
  
```

- else can only be used once at the end. The else clause gets executed if none of the above conditions are TRUE.
- The keywords else if and else must be in the same line as the closing brace }.

```

x <- 6
if (x > 0) {
    print("A positive number")
} else if (x < 0) {
    print("A negative number")
} else {
    print("It's zero!")
}

## [1] "A positive number"
  
```

Conditional statements can be written in various forms.

- Single-line statements don't need braces:

```
x <- 6
if (x > 0) y <- "Positive" else y <- "Not positive"
y
## [1] "Positive"
```

- The statement can return a value:

```
x <- 6
y <- if (x > 0) "Positive" else "Not positive"
y
## [1] "Positive"
```

[Task 1: if else Statement]

Inspect the outputs of the two code snippets below and explain how they work.

```
x <- 1:10
if (length(x)) "not empty" else "empty"
## [1] "not empty"

x <- numeric()
if (length(x)) "not empty" else "empty"
## [1] "empty"
```

[End of Task 1]

3.2.4 switch(): Select One from Several Alternatives

Syntax:

```
switch(EXPR, ...)
```

- switch(EXPR, ...) evaluates EXPR and chooses one of the alternatives accordingly.
- EXPR is the expression to evaluate.
- ... is the alternatives. If EXPR is a string, the alternatives will be named, perhaps except for one alternative to deal with the unmatched case.

```
cars$speed
## [1]  4  4  7  7  8  9 10 10 10 11 11 12 12 12 12 13 13 13 13 14 14 14 14
15 15
## [26] 15 16 16 17 17 17 18 18 18 18 19 19 19 20 20 20 20 20 22 23 24 24 24
24 25

type <- "MedianSpeed"
switch(type, MeanSpeed = mean(cars$speed), MedianSpeed = median(cars$speed),
"I don't understand")
```

```
## [1] 15
```

When calling `switch()`, you can leave the right hand side of `=` empty:

```
switch("b", a = 1, b = , c = 2, d = , e = 3, "nothing")
```

```
## [1] 2
```

```
switch("d", a = 1, b = , c = 2, d = , e = 3, "nothing")
```

```
## [1] 3
```

```
switch("f", a = 1, b = , c = 2, d = , e = 3, "nothing")
```

```
## [1] "nothing"
```

As you can see, when an empty result is found in a match, `switch()` will match the input to the next value that is not empty. This is an interesting feature that you can use if multiple inputs have the same output.

The above code snippet is equivalent to the following one:

```
switch("b", a = 1, b = 2, c = 2, d = 3, e = 3, "nothing")
```

```
## [1] 2
```

```
switch("d", a = 1, b = 2, c = 2, d = 3, e = 3, "nothing")
```

```
## [1] 3
```

```
switch("f", a = 1, b = 2, c = 2, d = 3, e = 3, "nothing")
```

```
## [1] "nothing"
```

[Task 2: `switch()`]

Rewrite the following code snippet with `switch()` (`stop()` stops execution and throws an error):

```
x = "a" # can be other values
```

```
if (x == "a") {
```

```
  "option 1"
```

```
} else if (x == "b") {
```

```
  "option 2"
```

```
} else if (x == "c") {
```

```
  "option 3"
```

```
} else {
```

```
  stop("Invalid x value")
```

```
}
```

```
## [1] "option 1"
```

[End of Task 2]

3.2.5 ifelse(): Conditional Element Selection

The function `ifelse()` performs **conditional element selection**.

Syntax:

```
ifelse(test, yes, no)
```

- `test` is a logical vector of conditions, `yes` and `no` are two vectors we want to select elements from.
- `ifelse(test, yes, no)` returns a vector which is filled with elements selected from either `yes` or `no`, depending on whether the element of `test` is TRUE or FALSE.
- If `yes` or `no` are too short, their elements are **recycled**.

```
set.seed(0)
x <- rnorm(5, 0, 2)
ifelse(x > 0, x, -x)  # `x > 0` is a logical vector of conditions
## [1] 2.5259086 0.6524667 2.6595985 2.5448586 0.8292829
ifelse(x > 0, x, 0)   # the `no` vector (0) is recycled
## [1] 2.5259086 0.0000000 2.6595985 2.5448586 0.8292829
```

3.3 Iterative Statements

Iterative statements (loops) allow us to execute a set of statements repeatedly.

3.3.1 for: Definite Loops

`for` is used for **definite loops** when the number of times to repeat is clear in advance.

Syntax:

```
for (counter in sequence){
  statements
}
```

- `counter` is a variable that iterates over the sequence (or a vector or a list).
- `statements` (called the *body* of the loop) are repeatedly run through iterations.

```
x <- c("a", "b", "c", "d")
for (i in 1:4) print(x[i])

## [1] "a"
## [1] "b"
## [1] "c"
## [1] "d"

for (letter in x) print(letter)
```

```
## [1] "a"
## [1] "b"
## [1] "c"
## [1] "d"
```

for loops can be nested (arbitrarily deeply).

```
x <- matrix(1:6, 2, 3)
x

##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6

for (i in seq_len(nrow(x))) {
  for (j in seq_len(ncol(x))) {
    print(x[i, j])
  }
}

## [1] 1
## [1] 3
## [1] 5
## [1] 2
## [1] 4
## [1] 6
```

Here, `seq_len()` is very useful in loops. Another useful function is `seq_along()`.

```
x <- c(1, 50, 20, 12)
seq_len(length(x))

## [1] 1 2 3 4

seq_along(x)

## [1] 1 2 3 4

1:length(x)

## [1] 1 2 3 4

# for a vector with length >= 1, the above three give the same result
# but for a vector of length 0, there is a difference:
x <- c()
seq_len(length(x))

## integer(0)

seq_along(x)

## integer(0)

1:length(x)    # returns a vector of length 2 in decreasing order
```



```
## [1] 1 0
```

[Task 3: Grade Assignment]

Run the following code snippet to generate a sequence of random numbers.

- The function `runif(n, min = 0, max = 1)` draws random numbers from a uniform distribution (?runif to learn more).
- The function `round()` rounds the values in its first argument to the specified number of decimal places (?round to learn more).
- `set.seed(0)` is used here to ensure the same set of random numbers are generated no matter when and where you run this code (otherwise, `runif()` is likely to generate different random numbers across different calls).

```
set.seed(0)
scores <- round(runif(50, 45, 100), 1)
scores

## [1] 94.3 59.6 65.5 76.5 95.0 56.1 94.4 97.0 81.3 79.6 48.4 56.3 54.7 82.8
## [16] 66.1 87.3 72.4 84.5 99.6 65.9 87.8 96.4 56.7 80.8 51.9 59.7 66.2 45.7 66.0
## [31] 92.8 63.7 71.5 78.0 72.1 55.2 90.5 81.8 88.7 50.9 84.8 67.6 90.2 80.6 88.1
## [46] 75.4 74.1 88.4 46.3 71.2 85.3
```

(a) Suppose that these numbers represent students' scores of a course. Write a **for loop** to convert these numbers to letter grades. The letter grade will be "A" if a score is at least 90. The letter grade will be "B" if a score is between 80 and 90. The letter grade will be "F" if a score is lower than "80". Create a variable called `grades` to store the result.

Tips:

- Create an empty character vector to fill things up later;
- A for loop is needed to traverse these 50 scores;
- The if statement may need to use its `else if` and `else` clause.

(b) Use `ifelse()` to do the same thing with a single line of code.

[End of Task 3]

3.3.2 while: Conditional Loops

`while` is used for **conditional loops** when we know the condition to stop the iterations.




Fig. 4 Flowchart of while Loops

Syntax:

```
while (condition) {
  statements
}
```

- R will run the statements repeatedly while the condition is TRUE.
- In other words, R will run the statements repeatedly until the condition turns FALSE.

```
i <- 1
x <- c("a", "b", "c", "d")
while (i <= length(x)) {
  print(x[i])
  i <- i+1
}

## [1] "a"
## [1] "b"
## [1] "c"
## [1] "d"
```

Compared with for loops, while loops are more general, in that every for loop could be replaced with a while loop (but not vice versa).

[Task 4: Least Common Multiple]

The least common multiple of a set of integers is the smallest number that is divisible by these integers. For example, the least common multiple of 2, 3 and 4 is 12.

Suppose that an arbitrary number of values are stored in a vector named `find.lcm`. Use conditionals and/or loops to find the least common multiple for values in this vector.

For example, for `find.lcm <- c(4, 5, 7)`, this code should yield 140; for `find.lcm <- c(3, 5, 7, 23)`, this code should return 2415; and `find.lcm <- 2:14`, this code should yield 360360.

[End of Task 4]

3.3.3 repeat: Infinite Loops

Syntax:

```
repeat {  
  statements  
}
```

- statements will be executed repeatedly until we press the ESCAPE button or quit R.
- We can break out of a repeat loop by including a break statement via conditionals. Then repeat does the same thing as while does.
- Infinite loops should generally be *avoided*, even if they are theoretically correct.

```
i <- 1  
x <- c("a", "b", "c", "d")  
repeat {  
  if (i > length(x)) # try: delete these two lines and run  
    break  
  print(x[i])  
  i<-i+1  
}  
  
## [1] "a"  
## [1] "b"  
## [1] "c"  
## [1] "d"
```

3.3.4 next and break: Control Loops Explicitly

- next skips the rest of the current iteration and starts the next iteration.

```
for(i in 1:8) {  
  if(i <= 5) next # skip the first 5 iterations  
  print(i)  
}  
  
## [1] 6  
## [1] 7  
## [1] 8  
  
print(paste("no. of iterations:", i))  
## [1] "no. of iterations: 8"
```

- break terminates the loop immediately.

```

for(i in 1:8) {
  if(i >= 5) break  # jump out of the loop when i >= 5
  print(i)
}

## [1] 1
## [1] 2
## [1] 3
## [1] 4

print(paste("no. of iterations:", i))

## [1] "no. of iterations: 5"

```

3.3.5 Vectorized Operations

Vectorized operations are always preferred to explicit loops when implementing element-wise operations, and should be leveraged whenever possible.

```

a = 1:5
b = 11:15

# vectorization:
c <- a + b
c

## [1] 12 14 16 18 20

# explicit loop:
c <- vector("numeric", length(a))
for (i in seq_along(a)) c[i] <- a[i] + b[i]
c

## [1] 12 14 16 18 20

```

Benefits of vectorized operations:

- Clarity: the syntax is about what we're doing.
- Abstraction: the syntax hides how the computer does it.
- Concision: we write less.
- Generality: same syntax works for numbers, vectors, arrays, ...
- Speed: modifying big vectors over and over is slow in R; work gets done by optimized low-level code.

3.4 Summary

Control structures allow us to control the flow of an R program.

Common control structures in R:

- **Conditional Statements (Conditionals):** execute different sets of statements, depending on whether a condition is true or false.
 - `if/else`: test a condition and act on it
 - `switch()`: select one from several alternatives
 - `ifelse()`: conditional element selection
- **Iterative Statements (Loops):** execute a set of statements repeatedly.
 - `for`: execute a loop for a fixed number of times
 - `while`: repeat a loop while a condition is true
 - `repeat`: run an infinite loop (must break out of it to stop)
 - `next`: skip the current iteration of a loop;
 - `break`: break out of a loop

[Task 5: Groupwise data summarization]

We're going to look at a data set on 97 men who have prostate cancer (from the book [The Elements of Statistical Learning](#)).

Download the data **pros.dat** from Canvas into the directory where you put your in-class exercise R script.

In this data, there are 9 variables measured on these 97 men:

- `lpsa`: log PSA score
- `lcavol`: log cancer volume
- `lweight`: log prostate weight
- `age`: age of patient
- `lbph`: log of the amount of benign prostatic hyperplasia
- `svi`: seminal vesicle invasion
- `lcp`: log of capsular penetration
- `gleason`: Gleason score
- `pgg45`: percent of Gleason scores 4 or 5

Use the `read.table()` function to load the data into your R session, and store it as a data frame `pros.dat`.

Before reading the data, you may need to set your *working directory* to the directory where you put the data and in-class exercise R script.

```
pros.dat <- read.table("pros.dat")
str(pros.dat)

## 'data.frame':    97 obs. of  9 variables:
## $ lcavol : num  -0.58 -0.994 -0.511 -1.204 0.751 ...
## $ lweight: num  2.77 3.32 2.69 3.28 3.43 ...
## $ age : int  50 58 74 58 62 50 64 58 47 63 ...
## $ lbph : num  -1.39 -1.39 -1.39 -1.39 -1.39 ...
## $ svi : int  0 0 0 0 0 0 0 0 0 0 ...
## $ lcp : num  -1.39 -1.39 -1.39 -1.39 -1.39 ...
## $ gleason: int  6 6 7 6 6 6 6 6 6 6 ...
## $ pgg45 : int  0 0 20 0 0 0 0 0 0 0 ...
## $ lpsa : num  -0.431 -0.163 -0.163 -0.163 0.372 ...
```

```
head(pros.dat)
```

```
##      lcavol  lweight age      lbph svi      lcp gleason pgg45      lpsa
## 1 -0.5798185 2.769459 50 -1.386294 0 -1.386294      6      0 -0.4307829
## 2 -0.9942523 3.319626 58 -1.386294 0 -1.386294      6      0 -0.1625189
## 3 -0.5108256 2.691243 74 -1.386294 0 -1.386294      7     20 -0.1625189
## 4 -1.2039728 3.282789 58 -1.386294 0 -1.386294      6      0 -0.1625189
## 5  0.7514161 3.432373 62 -1.386294 0 -1.386294      6      0  0.3715636
## 6 -1.0498221 3.228826 50 -1.386294 0 -1.386294      6      0  0.7654678
```

We randomly assign the 97 observations into 10 groups using the function `sample()`:

```
set.seed(0)
grouping <- sample(1:10, size = 97, replace = T)
grouping  # each number represents the group the corresponding observation
is assigned to

## [1] 9 4 7 1 2 7 2 3 1 5 5 10 6 10 7 9 5 5 9 9 5 5 2
## [26] 1 4 3 6 10 10 6 4 4 10 9 7 6 9 8 9 7 8 6 10 7 3 10
## [51] 2 2 6 6 1 3 3 8 6 7 6 8 7 1 4 8 9 9 7 4 7 6 1
## [76] 1 9 7 7 3 6 2 10 10 7 3 2 10 1 10 10 8 10 5 7 8 5
```

`sample()` generates a sample of a fixed size (i.e., 97) from 1:10 by sampling with replacement. (?sample for more information).

(a) Write a for loop (possibly nested) to calculate the groupwise sum of each column. Store the result in a **10 by 9** (10 groups and 9 columns) matrix called `groupwise.sum`. Your code should be as concise as possible.

Tips:

- You may need a nested loop with the outer loop iterating over columns and the inner loop iterating over groups;
- Use `grouping` with Boolean indexing to find observations assigned into one group;

- Use matrix indexing to fill results in `groupwise.sum`.

(b) Convert the `pros.dat` data frame to a matrix called `pros.mat` using the following code:

```
pros.mat <- as.matrix(pros.dat)
```

Then write a for loop to do the same thing. Now you only need a flat loop with `colSums()`.

(c) Figure out how to use a function named `rowsum()` with `pros.mat` to achieve the same goal.

[End of Task 5]