

ISOM 3390: Business Programming in R

Instructor: Bingjie Qian (PhD candidate)

Summer 2020, HKUST

Topic 6: Data Wrangling: A Tidy Data Approach

In the first 5 topics, we have learned the basic skills of R programming. Now we are ready to use R to tackle data science and business analytics problems.

6.1 An Overview of Tidy Data

In practice, there are multiple ways to store and present a data set, but they are not equally useful.

Suppose we have the data on 5 songs and their weekly ranks in the "Billboard Top 100" over 3 weeks. We can store the data in the *wide format* or in the *long format* as follows.

track	date.entered	week1	week2	week3
What A Girl Wants	11/27/1999	71	51	28
With Arms Wide Open	5/13/2000	84	78	76
Try Again	3/18/2000	59	53	38
Thank God I Found You	12/11/1999	82	68	50
Breathe	11/6/1999	81	68	62

Table 1. Wide Format

track	date.entered	week	position
What A Girl Wants	11/27/1999	1	71
What A Girl Wants	11/27/1999	2	51
What A Girl Wants	11/27/1999	3	28
With Arms Wide Open	5/13/2000	1	84
With Arms Wide Open	5/13/2000	2	78
With Arms Wide Open	5/13/2000	3	76
Try Again	3/18/2000	1	59
Try Again	3/18/2000	2	53
Try Again	3/18/2000	3	38
Thank God I Found You	12/11/1999	1	82
Thank God I Found You	12/11/1999	2	68
Thank God I Found You	12/11/1999	3	50
Breathe	11/6/1999	1	81
Breathe	11/6/1999	2	68
Breathe	11/6/1999	3	62

Table 2. Long Format

- In the wide format, each row represents *a specific song*.
- In the long format, each row represents *a specific song in a specific week*.

If we want to generate a plot that shows each song's ranking variation from the first to the last week, we need to use week number as the variable for x-axis. It is easy to do so with the long format, but with the wide format, week number is not a variable.

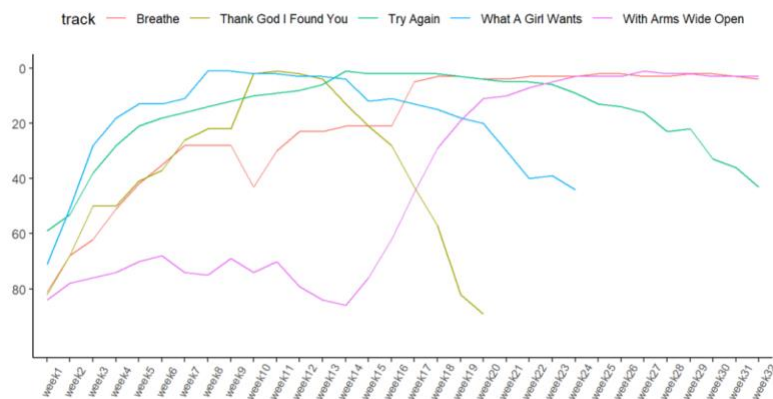


Fig 1. Ranking Plot

Besides, if some songs did not stay in the Billboard for all the weeks, you may find a lot of missing values (NA) with the wide format.

The dataset `billboard.csv` (download from Canvas) records weekly ranks of 317 songs appearing in the Billboard Top 100 over a certain time period.

```
billboard <- read.csv("billboard.csv")
head(billboard) # wide format
```

[illegible]

NA											
## 3	37	46	47	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 4	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 5	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 6	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
##	week42	week43	week44	week45	week46	week47	week48	week49	week50	week51	
week52											
## 1	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 2	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 3	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 4	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 5	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 6	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
##	week53	week54	week55	week56	week57	week58	week59	week60	week61	week62	
week63											
## 1	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 2	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 3	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 4	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 5	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
## 6	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA	NA
NA											
##	week64										
## 1	NA										
## 2	NA										
## 3	NA										
## 4	NA										
## 5	NA										
## 6	NA										

Common Causes of Data Messiness

Some common causes of messiness in data:

- Values are stored in column headers. E.g., in `billboard`, column headers are `week1`, `week2`, ..., `week64`, which stores the values of week number (1 to 64).
- Multiple types of observations in one table. E.g., the `billboard` dataset contains two types of observations: the song's meta information (`artist`, `time`, `genre`, `date.entered`, `date.peakd`) and its rank in each week.
- Multiple variables are stored in one column.
- One type in multiple tables.

"Tidy datasets are all alike, but every messy dataset is messy in its own way." - Hadley Wickham

Tidy data

Tidy data is a standard way of mapping the meaning of a dataset to its structure.

- The following three interrelated rules make a dataset tidy:
 - Each variable forms a column.
 - Each observation forms a row.
 - Each type of observation forms a table.
- Two main advantages of tidy data:
 - Enforcing a consistent data structure makes it easier to learn the tools that work with it because of the **underlying uniformity**.
 - Placing variables in columns allows R's **vectorized** nature to shine.
- A popular set of "tidy" tools known as the **tidyverse** allows us to transition smoothly from different stages of data analysis.

tidyverse: R Packages for Data Science

tidyverse is a collection of packages that share an underlying design philosophy, grammar and data structures, and are designed to work together naturally.

Install the complete tidyverse packages with:

```
# install.packages("tidyverse")
library(tidyverse)

## — Attaching packages ————— tidyverse
## 1.3.0 —

## ✓ ggplot2 3.3.2      ✓ purrr  0.3.4
## ✓ tibble  3.0.1      ✓ dplyr  1.0.0
```

```
## ✓ tidyr 1.1.0 ✓ stringr 1.4.0
## ✓ readr 1.3.1 ✓ forcats 0.5.0

## — Conflicts —————
tidyverse_conflicts() —
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
```

`install.packages("tidyverse")` installs the complete collection of packages in tidyverse.

`library(tidyverse)` loads the core tidyverse packages: `ggplot2`, `dplyr`, `tidyr`, `readr`, `purrr`, `tibble`, `stringr` and `forcats`.

Other packages with more specialized usage in tidyverse (e.g., `rvest`, `lubridate`, etc.) need to be loaded with their own `library()` calls.

6.2 Tibbles

One of the unifying features of tidyverse is a data structure called **tibble**.

Tibbles are a modern remake of the `data.frame` and encapsulate best practices for data frames.

Use `tibble()` to create a new tibble from column vectors:

```
tb <- tibble(x = 1:5, y = 1, z = x^2 + y) # 1 is recycled
tb

## # A tibble: 5 x 3
##       x     y     z
##   <int> <dbl> <dbl>
## 1     1     1     2
## 2     2     1     5
## 3     3     1    10
## 4     4     1    17
## 5     5     1    26

str(tb)

## tibble [5 × 3] (S3: tbl_df/tbl/data.frame)
## $ x: int [1:5] 1 2 3 4 5
## $ y: num [1:5] 1 1 1 1 1
## $ z: num [1:5] 2 5 10 17 26
```

It evaluates the arguments *sequentially* so that we can refer to variables just created, i.e., refer to `x` and `y` when creating `z`.

Differences between Tibbles and Data Frames

- It has a refined print method. It shows only the first 10 rows, fits all the columns on screen, and reports the type of each column along the name.

```
billboard_tbl <- as_tibble(billboard)
billboard_tbl

## # A tibble: 317 x 70
##   artist track time genre date.entered date.peaked week1 week2 week3
##   <chr> <chr> <chr> <chr> <chr> <chr> <int> <int> <int>
##   <int>
## 1 Desti... Inde... 3:38 Rock 9/23/2000 11/18/2000 78 63 49
## 33
## 2 Santa... Mari... 4:18 Rock 2/12/2000 4/8/2000 15 8 6
## 5
## 3 Savag... I Kn... 4:07 Rock 10/23/1999 1/29/2000 71 48 43
## 31
## 4 Madon... Music 3:45 Rock 8/12/2000 9/16/2000 41 23 18
## 14
## 5 Aguil... Come... 3:38 Rock 8/5/2000 10/14/2000 57 47 45
## 29
## 6 Janet Does... 4:17 Rock 6/17/2000 8/26/2000 59 52 43
## 30
## 7 Desti... Say ... 4:31 Rock 12/25/1999 3/18/2000 83 83 44
## 38
## 8 Igles... Be W... 3:36 Latin 4/1/2000 6/24/2000 63 45 34
## 23
## 9 Sisco Inco... 3:52 Rock 6/24/2000 8/12/2000 77 66 61
## 61
## 10 Lones... Amaz... 4:25 Coun... 6/5/1999 3/4/2000 81 54 44
## 39
## # ... with 307 more rows, and 60 more variables: week5 <int>, week6 <int>,
## # week7 <int>, week8 <int>, week9 <int>, week10 <int>, week11 <int>,
## # week12 <int>, week13 <int>, week14 <int>, week15 <int>, week16 <int>,
## # week17 <int>, week18 <int>, week19 <int>, week20 <int>, week21 <int>,
## # week22 <int>, week23 <int>, week24 <int>, week25 <int>, week26 <int>,
## # week27 <int>, week28 <int>, week29 <int>, week30 <int>, week31 <int>,
## # week32 <int>, week33 <int>, week34 <int>, week35 <int>, week36 <int>,
## # week37 <int>, week38 <int>, week39 <int>, week40 <int>, week41 <int>,
## # week42 <int>, week43 <int>, week44 <int>, week45 <int>, week46 <int>,
## # week47 <int>, week48 <int>, week49 <int>, week50 <int>, week51 <int>,
## # week52 <int>, week53 <int>, week54 <int>, week55 <int>, week56 <int>,
## # week57 <int>, week58 <int>, week59 <int>, week60 <int>, week61 <int>,
## # week62 <int>, week63 <int>, week64 <int>
```

- It never changes an input's type. With data frames, strings are often coerced to factors, and you have to set `stringsAsFactors = FALSE` in order to avoid this kind of coercion.
- It never changes the names of variables.

```
names(data.frame("crazy name" = 1)) # data frame changes space to dot
```

```
## [1] "crazy.name"
names(tibble("crazy name" = 1))      # tibble does not change the names
## [1] "crazy name"
```

- It does not allow setting row.names. Column names can be set using names.

```
row.names(tb) <- letters[1:5]
```

```
## Warning: Setting row names on a tibble is deprecated.
```

- It clearly distinguishes between [and [(. With tibbles, [always returns another *tibble*.

```
# data frame:
```

```
billboard[1:3, "week1"]      # return a vector
```

```
## [1] 78 15 71
```

```
billboard[1:3, "week1", drop = FALSE]  # return a data frame
```

```
##   week1
```

```
## 1    78
```

```
## 2    15
```

```
## 3    71
```

```
# tibble:
```

```
billboard_tbl[1:3, "week1"]  # return a tibble
```

```
## # A tibble: 3 x 1
```

```
##   week1
```

```
##   <int>
```

```
## 1    78
```

```
## 2    15
```

```
## 3    71
```

Remark: In the subsequent lectures, we will use the term *tibble* and *data frame* interchangeably when referring to tibbles.

6.3 tidyr for Data Tidying

tidyr functions help us *get data into the tidy data format*.

Two fundamental operations for data tidying:

- Converting data between wide and long formats.
 - pivot_longer() makes "wide" data longer.
 - pivot_wider() makes "long" data wider.
- Splitting and combining character columns.

- `separate()` splits a single column (that represents multiple variables) into multiple columns.
- `unite()` combines multiple columns into a single column.

6.3.1 `pivot_longer()`: Make a Wide Data Longer

In the following example, we will use a simpler dataset `billboard2.csv`, which only has 5 rows and 5 columns.

```
billboard <- read.csv("billboard2.csv")
billboard_tbl <- as_tibble(billboard)
billboard_tbl # wide format
```

```
## # A tibble: 5 x 5
```

	track	date.entered	week1	week2	week3
	<chr>	<chr>	<int>	<int>	<int>
## 1	What A Girl Wants	11/27/1999	71	51	28
## 2	With Arms Wide Open	5/13/2000	84	78	76
## 3	Try Again	3/18/2000	59	53	38
## 4	Thank God I Found You	12/11/1999	82	68	50
## 5	Breathe	11/6/1999	81	68	62

This data set is stored in the *wide* format.

Use the function `pivot_longer()` to turn it into a *long* format:

```
billboard_tbl_l <- pivot_longer(billboard_tbl, cols = week1:week3, names_to =
"week", values_to = "position")
billboard_tbl_l # Long format
```

```
## # A tibble: 15 x 4
```

	track	date.entered	week	position
	<chr>	<chr>	<chr>	<int>
## 1	What A Girl Wants	11/27/1999	week1	71
## 2	What A Girl Wants	11/27/1999	week2	51
## 3	What A Girl Wants	11/27/1999	week3	28
## 4	With Arms Wide Open	5/13/2000	week1	84
## 5	With Arms Wide Open	5/13/2000	week2	78
## 6	With Arms Wide Open	5/13/2000	week3	76
## 7	Try Again	3/18/2000	week1	59
## 8	Try Again	3/18/2000	week2	53
## 9	Try Again	3/18/2000	week3	38
## 10	Thank God I Found You	12/11/1999	week1	82
## 11	Thank God I Found You	12/11/1999	week2	68
## 12	Thank God I Found You	12/11/1999	week3	50
## 13	Breathe	11/6/1999	week1	81
## 14	Breathe	11/6/1999	week2	68
## 15	Breathe	11/6/1999	week3	62

Usage:

```

pivot_longer(
  data,
  cols,
  names_to = "name",
  values_to = "value",
  values_drop_na = FALSE,
  ...
)

```

- `pivot_longer()` lengthens data, increasing the number of rows and decreasing the number of columns.
- `data`: the data frame to pivot.
- `cols`: the columns to pivot into longer format. The specification syntax is flexible, e.g., select all columns between `x` and `z` with `x:z` (inclusive), exclude `y` with `-y`.
- `names_to`: the name of the new column(s) created from the data stored in the column names of data.
- `values_to`: the name of the new column created from the data stored as the cell values of data.
- `values_drop_na`: whether missing values to be removed.

6.3.2 `pivot_wider()`: Make a Long Data Wider

Use `pivot_wider()` to turn `billboard_tbl_1` (*long* format) into a *wide* format:

```

billboard_tbl_1

## # A tibble: 15 x 4
##   track                date.entered week  position
##   <chr>                <chr>      <chr>    <int>
## 1 What A Girl Wants    11/27/1999 week1      71
## 2 What A Girl Wants    11/27/1999 week2      51
## 3 What A Girl Wants    11/27/1999 week3      28
## 4 With Arms Wide Open  5/13/2000  week1      84
## 5 With Arms Wide Open  5/13/2000  week2      78
## 6 With Arms Wide Open  5/13/2000  week3      76
## 7 Try Again            3/18/2000  week1      59
## 8 Try Again            3/18/2000  week2      53
## 9 Try Again            3/18/2000  week3      38
## 10 Thank God I Found You 12/11/1999 week1      82
## 11 Thank God I Found You 12/11/1999 week2      68
## 12 Thank God I Found You 12/11/1999 week3      50
## 13 Breathe             11/6/1999  week1      81
## 14 Breathe             11/6/1999  week2      68
## 15 Breathe             11/6/1999  week3      62

```

```
billboard_tbl_w <- pivot_wider(billboard_tbl_l, names_from = week,
values_from = position)
billboard_tbl_w
```

```
## # A tibble: 5 x 5
##   track          date.entered week1 week2 week3
##   <chr>          <chr>          <int> <int> <int>
## 1 What A Girl Wants 11/27/1999         71    51    28
## 2 With Arms Wide Open 5/13/2000         84    78    76
## 3 Try Again         3/18/2000         59    53    38
## 4 Thank God I Found You 12/11/1999        82    68    50
## 5 Breathe          11/6/1999         81    68    62
```

Usage:

```
pivot_wider(
  data,
  names_from = name,
  values_from = value,
  values_fill = NULL,
  ...
)
```

- `pivot_wider()` widens data, increasing the number of columns and decreasing the number of rows.
- `data`: a data frame to pivot.
- `names_from`: the column(s) to get the name of the output columns.
- `values_from`: the column(s) we want to use to populate the output columns.
- `values_fill`: what values to substitute if there are combinations that don't exist.

[Task 1: Converting Data between Wide and Long Formats]

Download the dataset `weeklyprice.csv` from Canvas, which contains time series data for stock prices for three large Internet companies (Facebook, Google, and Amazon) from Oct. 30, 2012 to Oct. 30, 2017.

Load the dataset to create a tibble named `weekly` using `read_csv()` from the `readr` package (the data import package in the `tidyverse`; return a tibble).

```
weekly <- read_csv("weeklyprice.csv", col_names = TRUE, col_types = cols())
weekly

## # A tibble: 3 x 263
##   company `2012-10-29` `2012-11-05` `2012-11-12` `2012-11-19` `2012-11-26`
##   <chr>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1 FB         21.2        19.2        23.6        24         28
```

```
## 2 GOOG          342.          329.          321.          332.          347.
## 3 AMZN           232.          226.          225.          240.          252.
## # ... with 257 more variables: `2012-12-03` <dbl>, `2012-12-10` <dbl>,
## #   `2012-12-17` <dbl>, `2012-12-24` <dbl>, `2012-12-31` <dbl>,
## #   `2013-01-07` <dbl>, `2013-01-14` <dbl>, `2013-01-21` <dbl>,
## #   `2013-01-28` <dbl>, `2013-02-04` <dbl>, `2013-02-11` <dbl>,
## #   `2013-02-18` <dbl>, `2013-02-25` <dbl>, `2013-03-04` <dbl>,
## #   `2013-03-11` <dbl>, `2013-03-18` <dbl>, `2013-03-25` <dbl>,
## #   `2013-04-01` <dbl>, `2013-04-08` <dbl>, `2013-04-15` <dbl>,
## #   `2013-04-22` <dbl>, `2013-04-29` <dbl>, `2013-05-06` <dbl>,
## #   `2013-05-13` <dbl>, `2013-05-20` <dbl>, `2013-05-27` <dbl>,
## #   `2013-06-03` <dbl>, `2013-06-10` <dbl>, `2013-06-17` <dbl>,
## #   `2013-06-24` <dbl>, `2013-07-01` <dbl>, `2013-07-08` <dbl>,
## #   `2013-07-15` <dbl>, `2013-07-22` <dbl>, `2013-07-29` <dbl>,
## #   `2013-08-05` <dbl>, `2013-08-12` <dbl>, `2013-08-19` <dbl>,
## #   `2013-08-26` <dbl>, `2013-09-02` <dbl>, `2013-09-09` <dbl>,
## #   `2013-09-16` <dbl>, `2013-09-23` <dbl>, `2013-09-30` <dbl>,
## #   `2013-10-07` <dbl>, `2013-10-14` <dbl>, `2013-10-21` <dbl>,
## #   `2013-10-28` <dbl>, `2013-11-04` <dbl>, `2013-11-11` <dbl>,
## #   `2013-11-18` <dbl>, `2013-11-25` <dbl>, `2013-12-02` <dbl>,
## #   `2013-12-09` <dbl>, `2013-12-16` <dbl>, `2013-12-23` <dbl>,
## #   `2013-12-30` <dbl>, `2014-01-06` <dbl>, `2014-01-13` <dbl>,
## #   `2014-01-20` <dbl>, `2014-01-27` <dbl>, `2014-02-03` <dbl>,
## #   `2014-02-10` <dbl>, `2014-02-17` <dbl>, `2014-02-24` <dbl>,
## #   `2014-03-03` <dbl>, `2014-03-10` <dbl>, `2014-03-17` <dbl>,
## #   `2014-03-24` <dbl>, `2014-03-31` <dbl>, `2014-04-07` <dbl>,
## #   `2014-04-14` <dbl>, `2014-04-21` <dbl>, `2014-04-28` <dbl>,
## #   `2014-05-05` <dbl>, `2014-05-12` <dbl>, `2014-05-19` <dbl>,
## #   `2014-05-26` <dbl>, `2014-06-02` <dbl>, `2014-06-09` <dbl>,
## #   `2014-06-16` <dbl>, `2014-06-23` <dbl>, `2014-06-30` <dbl>,
## #   `2014-07-07` <dbl>, `2014-07-14` <dbl>, `2014-07-21` <dbl>,
## #   `2014-07-28` <dbl>, `2014-08-04` <dbl>, `2014-08-11` <dbl>,
## #   `2014-08-18` <dbl>, `2014-08-25` <dbl>, `2014-09-01` <dbl>,
## #   `2014-09-08` <dbl>, `2014-09-15` <dbl>, `2014-09-22` <dbl>,
## #   `2014-09-29` <dbl>, `2014-10-06` <dbl>, `2014-10-13` <dbl>,
## #   `2014-10-20` <dbl>, `2014-10-27` <dbl>, ...
```

It is a wide format with 3 rows and 263 columns. Each row stores weekly closing prices for a stock during the aforementioned period.

(a) If we are interested in investigating the fluctuation of weekly stock prices, we need to transform the data from the wide format to the long format. Save the resulting tibble as `weekly_long`, which should have 3 columns, company, date, and price.

(b) Transform the data from the long format (`weekly_long`) back to the wide format. But this time, use company names as the column headers. The resulting tibble should have 4 columns, date, FB, GOOG, and AMZN.

[End of Task 1]

6.3.3 separate(): Separate a Character Column into Multiple Columns

Use the function `separate()` to separate the column `date.entered` in `billboard_tbl_1` into three columns `month`, `day` and `year`:

```
billboard_tbl_1  # one column: `date.entered`

## # A tibble: 15 x 4
##   track                date.entered week  position
##   <chr>                <chr>      <chr>    <int>
## 1 What A Girl Wants    11/27/1999 week1      71
## 2 What A Girl Wants    11/27/1999 week2      51
## 3 What A Girl Wants    11/27/1999 week3      28
## 4 With Arms Wide Open  5/13/2000  week1      84
## 5 With Arms Wide Open  5/13/2000  week2      78
## 6 With Arms Wide Open  5/13/2000  week3      76
## 7 Try Again            3/18/2000  week1      59
## 8 Try Again            3/18/2000  week2      53
## 9 Try Again            3/18/2000  week3      38
## 10 Thank God I Found You 12/11/1999 week1      82
## 11 Thank God I Found You 12/11/1999 week2      68
## 12 Thank God I Found You 12/11/1999 week3      50
## 13 Breathe             11/6/1999  week1      81
## 14 Breathe             11/6/1999  week2      68
## 15 Breathe             11/6/1999  week3      62

billboard_tbl_ls <- separate(billboard_tbl_1, col = date.entered, into =
c("month", "day", "year"))
billboard_tbl_ls  # three columns: `month`, `day` and `year`

## # A tibble: 15 x 6
##   track                month day  year  week  position
##   <chr>                <chr> <chr> <chr> <chr>    <int>
## 1 What A Girl Wants    11    27   1999 week1      71
## 2 What A Girl Wants    11    27   1999 week2      51
## 3 What A Girl Wants    11    27   1999 week3      28
## 4 With Arms Wide Open  5     13   2000 week1      84
## 5 With Arms Wide Open  5     13   2000 week2      78
## 6 With Arms Wide Open  5     13   2000 week3      76
## 7 Try Again            3     18   2000 week1      59
## 8 Try Again            3     18   2000 week2      53
## 9 Try Again            3     18   2000 week3      38
## 10 Thank God I Found You 12     11   1999 week1      82
## 11 Thank God I Found You 12     11   1999 week2      68
## 12 Thank God I Found You 12     11   1999 week3      50
## 13 Breathe             11     6   1999 week1      81
## 14 Breathe             11     6   1999 week2      68
## 15 Breathe             11     6   1999 week3      62
```

Usage:

```
separate(  
  data,  
  col,  
  into,  
  sep = "[^[:alnum:]]+",  
  ...  
)
```

- `col`: the column to be split apart.
- `into`: the names of new variables to create as character vectors.
- `sep`: the separator between columns.
 - specified either by a *regular expression* (more on it later) or by *position* (e.g., `sep = c(3, -5)`);
 - If character, `sep` is interpreted as a *regular expression*. The default value is a regular expression that matches *non-alphanumeric characters*.
 - If numeric, `sep` is interpreted as *character positions* to split at. Positive values start at 1 at the far-left of the string; negative value start at -1 at the far-right of the string.

6.3.4 `unite()`: Unite Multiple Columns into One by Pasting Strings Together

```
billboard_tbl_ls # three columns: `month`, `day` and `year`
```

```
## # A tibble: 15 x 6  
##   track                month day   year week  position  
##   <chr>              <chr> <chr> <chr> <chr>    <int>  
## 1 What A Girl Wants    11   27   1999 week1     71  
## 2 What A Girl Wants    11   27   1999 week2     51  
## 3 What A Girl Wants    11   27   1999 week3     28  
## 4 With Arms Wide Open  5    13   2000 week1     84  
## 5 With Arms Wide Open  5    13   2000 week2     78  
## 6 With Arms Wide Open  5    13   2000 week3     76  
## 7 Try Again            3    18   2000 week1     59  
## 8 Try Again            3    18   2000 week2     53  
## 9 Try Again            3    18   2000 week3     38  
## 10 Thank God I Found You 12    11   1999 week1     82  
## 11 Thank God I Found You 12    11   1999 week2     68  
## 12 Thank God I Found You 12    11   1999 week3     50  
## 13 Breathe             11    6    1999 week1     81  
## 14 Breathe             11    6    1999 week2     68  
## 15 Breathe             11    6    1999 week3     62
```

```
billboard_tbl_lu <- unite(billboard_tbl_ls, col = date.entered, day, month,  
  year, sep = "/")  
billboard_tbl_lu # one column: `date.entered`
```

```
## # A tibble: 15 x 4
##   track      date.entered week position
##   <chr>      <chr>      <chr>    <int>
## 1 What A Girl Wants 27/11/1999 week1     71
## 2 What A Girl Wants 27/11/1999 week2     51
## 3 What A Girl Wants 27/11/1999 week3     28
## 4 With Arms Wide Open 13/5/2000 week1     84
## 5 With Arms Wide Open 13/5/2000 week2     78
## 6 With Arms Wide Open 13/5/2000 week3     76
## 7 Try Again        18/3/2000 week1     59
## 8 Try Again        18/3/2000 week2     53
## 9 Try Again        18/3/2000 week3     38
## 10 Thank God I Found You 11/12/1999 week1     82
## 11 Thank God I Found You 11/12/1999 week2     68
## 12 Thank God I Found You 11/12/1999 week3     50
## 13 Breathe         6/11/1999 week1     81
## 14 Breathe         6/11/1999 week2     68
## 15 Breathe         6/11/1999 week3     62
```

Usage:

```
unite(data, col, ..., sep = "_", remove = TRUE, na.rm = FALSE)
```

- col: the name of the new column, as a string or symbol.
- ...: columns to unite.
- sep: separator to use between values.

[Task 2: Splitting and Combining Character Columns]

Continue the weekly stock prices problem in Task 1.

(a) Suppose we want to calculate monthly stock prices by averaging weekly prices. To do so, we need to first separate different parts of date values and store them as different variables. Work on `weekly_long` to get the following tibble and name it `monthly_long`:

```
# A tibble: 786 x 3
  company month  price
  <chr>   <chr>   <dbl>
1 FB      2012-10  21.2
2 FB      2012-11  19.2
3 FB      2012-11  23.6
4 FB      2012-11  24
5 FB      2012-11  28
6 FB      2012-12  27.5
7 FB      2012-12  26.8
8 FB      2012-12  26.3
9 FB      2012-12  25.9
```

```
10 FB      2012-12  28.8
# . with 776 more rows
```

Tips: split a date value by "-", drop the day part, and combine the year and month parts.

(b) Transform `monthly_long` from the long format to the wide format. call `pivot_wider()` again as you did in Task 1(b). But this time, include `values_fn = list(price = mean)` in the argument list.

[End of Task 2]

6.4 The Pipe Operator `%>%`

The `magrittr` package provides the pipe operator (`%>%`), which passes the result of one step (the left-hand side) as input for the next step (the right-hand side).

```
x %>% f      # equivalent to f(x)
```

First-argument rule:

```
x %>% f(y)    # equivalent to f(x, y)
```

The **argument placeholder** specifies where the piped input should land:

```
x %>% f(y, .)    # equivalent to f(y, x)
x %>% f(y, z=.)   # equivalent to f(y, z=x)
```

When the placeholder only appears in a nested expressions, the first-argument rule still applies:

```
x %>% f(y = nrow(.), z = ncol(.))    # equivalent to f(x, y = nrow(x), z = ncol(x))
```

In the above code, `x` is passed as the first argument of `f()`, and also to the place specified by `..`

The behavior can be overruled by enclosing the right-hand side in braces:

```
x %>% {f(y = nrow(.), z = ncol(.))}    # equivalent to f(y = nrow(x), z = ncol(x))
```

With `%>%`, we can process a data-object using *a sequence of operations* rather than *nested function calls*.

```
x %>% f %>% g %>% h    # equivalent to h(g(f(x)))
```

The intended aim of pipe operators is to *increase human readability* of written code.

It helps avoid nested function calls, minimize the need for local variables and function definitions, and add steps anywhere in the sequence of operations.

%>% is used throughout tidyverse and loaded automatically when using packages in it.

```
billboard_tbl %>% pivot_longer(cols = week1:week3, names_to = "week",
values_to = "position") %>% separate(date.entered, into = c("year", "month",
"day")) %>% subset(week == "week1") %>% .$position %>% mean

## [1] 75.4
```

Saving Pipelines

The input to the pipeline can itself be a placeholder:

```
num_unique <- . %>% unique %>% length
```

In this case, the pipeline describes a *function chain* that can be saved and re-used. It also has a different print method:

```
num_unique      # functional sequence

## Functional sequence with the following components:
##
## 1. unique(.)
## 2. length(.)
##
## Use 'functions' to extract the individual functions.
```

Apply the functional sequence to a data object:

```
x <- rep(1:5, times = c(2, 3, 2, 5, 2))
x %>% num_unique

## [1] 5

x %>% unique %>% length

## [1] 5

length(unique(x))

## [1] 5
```

[Task 3: Pipes to Base R]

For each of the following code blocks, which are written with pipes, write equivalent code in base R.

(a)

Pipes:

```
# `letters` is a built-in constant in R; ?letters
# `toupper` is a built-in function in R; ?toupper
```

```
# it is always good to build up your vocabulary of built-in functions
letters %>% toupper %>% paste(collapse="+")
```


```
## [1] "A+B+C+D+E+F+G+H+I+J+K+L+M+N+O+P+Q+R+S+T+U+V+W+X+Y+Z"
```

Base R:

(b)

Pipes:

```
rnorm(1000) %>% hist(breaks=30, main="N(0, 1) draws", col="pink", prob=TRUE)
```



Base R:

(c)

Pipes:

```
rnorm(1000) %>% hist(breaks=30, plot=FALSE) %>% .$density %>% max
```

```
## [1] 0.405
```

Base R:

[End of Task 3]

[Task 4: Base R to Pipes]

For each of the following code blocks, which are written in base R, write equivalent code with pipes (to do the same thing).

(a) Tips: use the dot . as seen in the lecture notes.

Base R:

```
paste("Your grade is", sample(c("A", "B", "C", "D", "F"), size = 1))  
## [1] "Your grade is C"
```

Pipes:

(b) Tips: use the dot . again, in order to index state.name directly in the last pipe command.

Base R:

```
# state.name and state.x77 are built-in datasets; ?state.x77; ?state.name;  
?which.max  
state.name[which.max(state.x77[, "Illiteracy"])]  
## [1] "Louisiana"
```

Pipes:

[End of Task 4]

6.5 dplyr for Data Manipulation

The package dplyr provides a consistent set of verbs that help streamline the data manipulation process.

- Single-table verbs:
 - select(data, variables): pick variables based on their names.
 - filter(data, conditions): pick observations based on conditions.
 - arrange(data, variables): reorder observations according to variables.
 - mutate(data, newvar = function): add new variables or transform existing variables.
 - summarize(data, newvar = function): collapse many values down to a single summary.

- `group_by(data, variables)`: group observations by variables.
- Two-table verbs: e.g., `inner_join(data1, data2, variables)`.
- Some other functions such as `slice()`, `rename()`, `transmute()`, `sample_n()` and `sample_frac()`, all of which we may find useful.

6.5.1 Select & Reorder Columns with `select()`

Let's first generate a data frame (tibble):

```
billboard <- as_tibble(read.csv("billboard.csv"))
billboard_l <- pivot_longer(billboard, starts_with("week"), names_to =
"week", values_to = "position", values_drop_na = TRUE)
billboard_l
```

```
## # A tibble: 5,306 x 8
##   artist      track      time  genre date.entered date.peaked week
##   <chr>      <chr>      <chr> <chr> <chr>      <chr>      <chr>
##   <int>
## 1 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week1
## 2 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week2
## 3 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week3
## 4 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week4
## 5 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week5
## 6 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week6
## 7 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week7
## 8 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week8
## 9 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week9
## 10 Destiny's ... Independent ... 3:38  Rock  9/23/2000    11/18/2000 week...
## # ... with 5,296 more rows
```

We can use `select()` to select (and optionally rename) variables in the data frame.

Usage:

```
select(.data, ...)
```

- `.data`: a data frame

- ...: one or more unquoted expressions separated by commas. Variable names can be used as if they were positions in the data frame, so expressions like `x:y` can be used to select a range of variables.

Example: From the tibble `billboard_1`, select the columns `track`, `week1` through `week3`, and `date.entered`, with the last being renamed to `date`:

```
billboard %>% select(track, week1:week3, date = date.entered)
```

```
## # A tibble: 317 x 5
```

	track	week1	week2	week3	date
	<chr>	<int>	<int>	<int>	<chr>
## 1	Independent Women Part I	78	63	49	9/23/2000
## 2	Maria, Maria	15	8	6	2/12/2000
## 3	I Knew I Loved You	71	48	43	10/23/1999
## 4	Music	41	23	18	8/12/2000
## 5	Come On Over Baby (All I Want Is You)	57	47	45	8/5/2000
## 6	Doesn't Really Matter	59	52	43	6/17/2000
## 7	Say My Name	83	83	44	12/25/1999
## 8	Be With You	63	45	34	4/1/2000
## 9	Incomplete	77	66	61	6/24/2000
## 10	Amazed	81	54	44	6/5/1999

```
## # ... with 307 more rows
```

Alternatively, we can replace `week1:week3` using a *selection helper* `num_range`:

```
billboard %>% select(track, num_range("week", 1:3), date = date.entered)
```

```
## # A tibble: 317 x 5
```

	track	week1	week2	week3	date
	<chr>	<int>	<int>	<int>	<chr>
## 1	Independent Women Part I	78	63	49	9/23/2000
## 2	Maria, Maria	15	8	6	2/12/2000
## 3	I Knew I Loved You	71	48	43	10/23/1999
## 4	Music	41	23	18	8/12/2000
## 5	Come On Over Baby (All I Want Is You)	57	47	45	8/5/2000
## 6	Doesn't Really Matter	59	52	43	6/17/2000
## 7	Say My Name	83	83	44	12/25/1999
## 8	Be With You	63	45	34	4/1/2000
## 9	Incomplete	77	66	61	6/24/2000
## 10	Amazed	81	54	44	6/5/1999

```
## # ... with 307 more rows
```

A variety of *helper functions* can be used to select variables based on their names.

- `starts_with("abc")` matches names that begin with "abc".
- `ends_with("abc")` matches names that end with "abc".
- `contains("abc")` matches names that contain "abc".

- `matches("(.)\\1")` selects variables that match a *regular expression* (more on this in Topic 9). This one matches any variables that contain repeated characters.
- `num_range("abc", 1:3)` matches `abc1`, `abc2` and `abc3`.
- `everything()`: all variables that haven't been specified.

```
billboard %>% select(track, date = date.entered, everything()) # reorder and
rename columns
```

```
## # A tibble: 317 x 70
##   track date   artist time  genre date.peaked week1 week2 week3 week4
week5
##   <chr> <chr> <chr>  <chr> <chr> <chr>          <int> <int> <int> <int>
<int>
##  1 Inde... 9/23... Desti... 3:38  Rock  11/18/2000      78    63    49    33
23
##  2 Mari... 2/12... Santa... 4:18  Rock   4/8/2000      15     8     6     5
2
##  3 I Kn... 10/2... Savag... 4:07  Rock   1/29/2000      71    48    43    31
20
##  4 Music 8/12... Madon... 3:45  Rock   9/16/2000      41    23    18    14
2
##  5 Come... 8/5/... Aguil... 3:38  Rock  10/14/2000      57    47    45    29
23
##  6 Does... 6/17... Janet   4:17  Rock   8/26/2000      59    52    43    30
29
##  7 Say ... 12/2... Desti... 4:31  Rock   3/18/2000      83    83    44    38
16
##  8 Be W... 4/1/... Igles... 3:36  Latin  6/24/2000      63    45    34    23
17
##  9 Inco... 6/24... Sisco   3:52  Rock   8/12/2000      77    66    61    61
61
## 10 Amaz... 6/5/... Lones... 4:25  Coun... 3/4/2000      81    54    44    39
38
## # ... with 307 more rows, and 59 more variables: week6 <int>, week7 <int>,
## #   week8 <int>, week9 <int>, week10 <int>, week11 <int>, week12 <int>,
## #   week13 <int>, week14 <int>, week15 <int>, week16 <int>, week17 <int>,
## #   week18 <int>, week19 <int>, week20 <int>, week21 <int>, week22 <int>,
## #   week23 <int>, week24 <int>, week25 <int>, week26 <int>, week27 <int>,
## #   week28 <int>, week29 <int>, week30 <int>, week31 <int>, week32 <int>,
## #   week33 <int>, week34 <int>, week35 <int>, week36 <int>, week37 <int>,
## #   week38 <int>, week39 <int>, week40 <int>, week41 <int>, week42 <int>,
## #   week43 <int>, week44 <int>, week45 <int>, week46 <int>, week47 <int>,
## #   week48 <int>, week49 <int>, week50 <int>, week51 <int>, week52 <int>,
## #   week53 <int>, week54 <int>, week55 <int>, week56 <int>, week57 <int>,
## #   week58 <int>, week59 <int>, week60 <int>, week61 <int>, week62 <int>,
## #   week63 <int>, week64 <int>
```

- `last_col()`: last variable, possibly with an offset.

```

billboard %>% select(track, last_col())

## # A tibble: 317 x 2
##   track                                week64
##   <chr>                                <int>
## 1 Independent Women Part I             NA
## 2 Maria, Maria                        NA
## 3 I Knew I Loved You                  NA
## 4 Music                               NA
## 5 Come On Over Baby (All I Want Is You) NA
## 6 Doesn't Really Matter               NA
## 7 Say My Name                         NA
## 8 Be With You                         NA
## 9 Incomplete                          NA
## 10 Amazed                             50
## # ... with 307 more rows

billboard %>% select(track, last_col(offset = 2)) # Set offset=n to select
the nth var from the end

## # A tibble: 317 x 2
##   track                                week62
##   <chr>                                <int>
## 1 Independent Women Part I             NA
## 2 Maria, Maria                        NA
## 3 I Knew I Loved You                  NA
## 4 Music                               NA
## 5 Come On Over Baby (All I Want Is You) NA
## 6 Doesn't Really Matter               NA
## 7 Say My Name                         NA
## 8 Be With You                         NA
## 9 Incomplete                          NA
## 10 Amazed                             42
## # ... with 307 more rows

```

6.5.2 Subset Rows with filter()

Usage:

```
filter(.data, ..., .preserve = FALSE)
```

- ...: Expressions that return a logical value, and are defined in terms of the variables in .data. If there are multiple expressions separated by commas, they are combined with the & operator.

Example: Look at the weekly ranks of a subset of songs (by the first 10 artists), and only look at the weeks when they rank higher than the average rank of all songs:

```

billboard_1 %>% filter(artist %in% unique(artist)[1:10] & position <
mean(position))

```

```
## # A tibble: 499 x 8
##   artist      track      time genre date.entered date.peaked week
position
##   <chr>      <chr>      <chr> <chr> <chr>      <chr>      <chr>
<int>
## 1 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week3
49
## 2 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week4
33
## 3 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week5
23
## 4 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week6
15
## 5 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week7
7
## 6 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week8
5
## 7 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week9
1
## 8 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week...
1
## 9 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week...
1
## 10 Destiny's ... Independent ... 3:38 Rock 9/23/2000 11/18/2000 week...
1
## # ... with 489 more rows
```

The operator `%in%` is a binary operator, indicating if there is a match or not for its left operand. `?'%in%'` for more information.

6.5.3 Rearrange Rows with `arrange()`

Usage:

```
arrange(.data, ..., .by_group = FALSE)
```

- `arrange()` orders the rows of a data frame by the values of selected columns.
- Use `desc()` to sort a variable in descending order.

Example: Sort weekly ranking positions first by artist, second by track, and last by position (descending order):

```
billboard_1 %>% arrange(artist, track, desc(position))

## # A tibble: 5,306 x 8
##   artist track      time genre date.entered date.peaked week
position
##   <chr>  <chr>      <chr> <chr> <chr>      <chr>      <chr>
<int>
## 1 2 Pac    Baby Don't Cry (... 4:22 Rap 2/26/2000 3/11/2000 week7
```



```

99
## 2 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week6
94
## 3 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week1
87
## 4 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week5
87
## 5 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week2
82
## 6 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week4
77
## 7 2 Pac    Baby Don't Cry (... 4:22 Rap    2/26/2000    3/11/2000    week3
72
## 8 2Ge+her The Hardest Part... 3:15 R&B    9/2/2000    9/9/2000    week3
92
## 9 2Ge+her The Hardest Part... 3:15 R&B    9/2/2000    9/9/2000    week1
91
## 10 2Ge+her The Hardest Part... 3:15 R&B    9/2/2000    9/9/2000    week2
87
## # ... with 5,296 more rows

```

6.5.4 Add, Modify, and Delete Columns with `mutate()` and `transmute()`

- `mutate()` adds new variables and *preserves* existing ones.
- `transmute()` adds new variables and *drops* existing ones.
- New variables overwrite existing variables of the same name.
- Variables can be removed by setting their value to NULL.

Example: Calculate changes in ranking positions between two adjacent weeks with `mutate()`:

```

billboard_1 %>% filter(track=="Maria, Maria") %>% .[c("track", "week",
"position")] %>% mutate(current = position, previous = lag(position), change
= previous - current, week = as.integer(str_extract(week, "\\d+")),
position = NULL)

## # A tibble: 26 x 5
##   track          week current previous change
##   <chr>         <int>   <int>   <int>   <int>
## 1 Maria, Maria     1     15      NA      NA
## 2 Maria, Maria     2      8     15       7
## 3 Maria, Maria     3      6      8       2
## 4 Maria, Maria     4      5      6       1
## 5 Maria, Maria     5      2      5       3
## 6 Maria, Maria     6      3      2      -1
## 7 Maria, Maria     7      2      3       1
## 8 Maria, Maria     8      2      2       0
## 9 Maria, Maria     9      1      2       1

```

```
## 10 Maria, Maria    10      1      1      0
## # ... with 16 more rows
```

Alternatively, we can use `transmute()`, no need to set `position = NULL`:

```
billboard_1 %>% filter(track=="Maria, Maria") %>% .[c("track", "week",
"position")] %>% transmute(current = position, previous = lag(position),
change = previous - position, week = as.integer(str_extract(week, "[\\d]+")))
```

```
## # A tibble: 26 x 4
##   current previous change week
##   <int>    <int>   <int> <int>
## 1     15      NA     NA     1
## 2      8     15      7     2
## 3      6      8      2     3
## 4      5      6      1     4
## 5      2      5      3     5
## 6      3      2     -1     6
## 7      2      3      1     7
## 8      2      2      0     8
## 9      1      2      1     9
## 10     1      1      0    10
## # ... with 16 more rows
```

In the above codes, `lag()` returns the *previous* values in a vector. Another useful function is `lead()` which returns the *next* values in a vector:

```
x <- 1:5
tibble(x_behind = lag(x), x, x_ahead = lead(x))
```

```
## # A tibble: 5 x 3
##   x_behind    x x_ahead
##   <int> <int>   <int>
## 1     NA     1     2
## 2      1     2     3
## 3      2     3     4
## 4      3     4     5
## 5      4     5    NA
```

`log()` and `lead()` are two examples of **window functions**.

A window function takes *n* inputs and returns *n* values. Examples include:

- Ranking and ordering functions: `row_number()`, `min_rank()`, `dense_rank()`, `cume_dist()`, `percent_rank()`, and `ntile()`.
- Offsets `lead()` and `lag()` allow us to compute differences and trends.
- Cumulative aggregates: `cumsum()`, `cummin()`, `cummax()` from base R; `cumall()`, `cumany()`, and `cummean()` from `dplyr`.

6.5.5 Group Data by group_by()

- Most data operations are done on *groups* defined by variables.
- `group_by()` takes an existing `tbl` and converts it into a grouped `tbl` where *operations are performed "by group"*.
- `ungroup()` removes grouping.

Example: Group weekly ranking positions first by week and then by artist:

```
billboard_1 %>% group_by(week, artist)

## # A tibble: 5,306 x 8
## # Groups:   week, artist [3,988]
##   artist      track      time genre date.entered date.peakd week
##   <chr>      <chr>      <chr> <chr> <chr>      <chr>      <chr>
##   <int>
## 1 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week1
## 78
## 2 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week2
## 63
## 3 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week3
## 49
## 4 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week4
## 33
## 5 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week5
## 23
## 6 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week6
## 15
## 7 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week7
## 7
## 8 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week8
## 5
## 9 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week9
## 1
## 10 Destiny's ... Independent ... 3:38  Rock  9/23/2000   11/18/2000 week...
## 1
## # ... with 5,296 more rows
```

Compared to the original `tbl`, the grouped `tbl` shows "Groups: week, artist [3,988]".

```
billboard_1 %>% group_by(week, artist) %>% groups      # return a list of
group names

## [[1]]
## week
##
## [[2]]
## artist
```

```
billboard_1 %>% group_by(week, artist) %>% group_vars # return a character
vector of group names

## [1] "week" "artist"
```

By default, group_by() overrides existing grouping:

```
billboard_1 %>% group_by(week, artist) %>% group_by(track) %>% group_vars

## [1] "track"
```

Set .add = TRUE to instead append grouping:

```
billboard_1 %>% group_by(week, artist) %>% group_by(track, .add = TRUE) %>%
group_vars

## [1] "week" "artist" "track"
```

Use ungroup() to removing grouping:

```
billboard_1 %>% group_by(week, artist) %>% ungroup %>% group_vars

## character(0)
```

[Task 5: Fastest 100m sprint times]

We read in a data frame sprint.m.df containing of the fastest times ever recorded for the 100m sprint in men's track.

```
sprint.m.df <- read_tsv("sprint.m.dat", col_names = TRUE)

## Parsed with column specification:
## cols(
##   Rank = col_double(),
##   Time = col_double(),
##   Wind = col_double(),
##   Name = col_character(),
##   Country = col_character(),
##   Birthdate = col_character(),
##   City = col_character(),
##   Date = col_character()
## )

sprint.m.df

## # A tibble: 2,988 x 8
##   Rank Time Wind Name Country Birthdate City Date
##   <dbl> <dbl> <dbl> <chr> <chr> <chr> <chr> <chr>
## 1     1  9.58  0.9 Usain Bolt JAM 21.08.86 Berlin
16.08.2009
## 2     2  9.63  1.5 Usain Bolt JAM 21.08.86 London
```

```

05.08.2012
## 3      3  9.69    0   Usain Bolt    JAM    21.08.86   Beijing
16.08.2008
## 4      3  9.69    2    Tyson Gay    USA    09.08.82   Shanghai
20.09.2009
## 5      3  9.69  -0.1 Yohan Blake    JAM    26.12.89   Lausanne
23.08.2012
## 6      6  9.71    0.9 Tyson Gay    USA    09.08.82   Berlin
16.08.2009
## 7      7  9.72    1.7 Usain Bolt    JAM    21.08.86   New York City
31.05.2008
## 8      7  9.72    0.2 Asafa Powell  JAM    23.11.82   Lausanne
02.09.2008
## 9      9  9.74    1.7 Asafa Powell  JAM    23.11.82   Rieti
09.09.2007
## 10     9  9.74    0.9 Justin Gatlin USA    10.02.82   Ad-Dawah
15.05.2015
## # ... with 2,978 more rows

```

Compute, for each country, the quadruple: (Name, City, Country, Time) corresponding to the athlete with the fastest time among athletes from that country, and display the first 10 results ordered by increasing time. If there are ties, then show all the results that correspond to the fastest time.

Tips: group by Country; find all the results with the fastest time for each country; and arrange the results by increasing time.

[End of Task 5]

6.5.6 Split-and-Apply Data Analysis with `group_by()` and `summarise()`

Form a summary table showing the number of weeks on chart, average ranking, and peak position of *each track*:

```

billboard_1 %>% group_by(track) %>% summarise("weeks on chart" = n(),
"average position" = mean(position), "peak position" = min(position))

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 316 x 4
##   track                                `weeks on chart` `average positio...` `peak
position`
##   <chr>                                <int>          <dbl>
<int>
## 1 (Hot S**t) Country Grammar          34            30.9
7
## 2 3 Little Words                      9            94.4
89
## 3 911                                19            60

```

```

38
## 4 A Country Boy Can Survive          3          86.7
75
## 5 A Little Gasoline                  6          89.8
75
## 6 A Puro Dolor (Purest Of P...      26          49.5
26
## 7 Aaron's Party (Come Get I...      15          66.3
35
## 8 Absolutely (Story Of A Gi...      27          26.5
6
## 9 All Good?                          3          97.3
96
## 10 All The Small Things              23          33.3
6
## # ... with 306 more rows

```

Calculate *every artist's* best *weekly* ranking (an artist may have multiple tracks):

```

billboard_1 %>% group_by(week, artist) %>% summarise("peak position" =
min(position))

## `summarise()` regrouping output by 'week' (override with `.groups`
argument)

## # A tibble: 3,988 x 3
## # Groups:   week [64]
##   week artist          `peak position`
##   <chr> <chr>              <int>
## 1 week1 2 Pac              87
## 2 week1 2Ge+her            91
## 3 week1 3 Doors Down       76
## 4 week1 504 Boyz           57
## 5 week1 98?                51
## 6 week1 A*Teens            97
## 7 week1 Aaliyah            59
## 8 week1 Adams, Yolanda      76
## 9 week1 Adkins, Trace       84
## 10 week1 Aguilera, Christina 50
## # ... with 3,978 more rows

```

Each call to `summarise()` removes a layer of grouping.

```

billboard_1 %>% group_by(week, artist) %>% summarise("peak position" =
min(position)) %>% group_vars()

## `summarise()` regrouping output by 'week' (override with `.groups`
argument)

## [1] "week"

```




Fig 2. Group and Summarise

In the above codes, `n()`, `mean()`, and `min()` are examples of **aggregation functions**.

An aggregation function takes `n` inputs and return a single value. Examples include:

- `min()`, `max()`, `...`, `median()`, etc. from base R.
- `n()`: the number of observations in the current group.
- `n_distinct(x)`: the number of unique values in `x`.
- `first(x)`, `last(x)` and `nth(x, n)`
 - work similarly to `x[1]`, `x[length(x)]`, and `x[n]` but give more control over the result.

6.5.7 By-Group Operations

All dplyr verbs discussed so far can combine naturally with `group_by()` to enable "by group" operations.

*** Grouped `arrange()`***

By default, it ignores groups:

```
billboard_1 %>% group_by(track) %>% .[c("artist", "track", "week",
"position")] %>% arrange(desc(position))
```

```
## # A tibble: 5,306 x 4
## # Groups:   track [316]
##   artist                track                week  position
##   <chr>                 <chr>                <chr>    <int>
## 1 Nelly                 (Hot S**t) Country Grammar week1      100
## 2 Martin, Ricky        She Bangs                week18     100
## 3 Lil Bow Wow           Bounce With Me           week22     100
## 4 Spears, Britney       Lucky                    week11     100
## 5 Hoku                  Another Dumb Blonde       week14     100
```

```
## 6 Sugar Ray           Falls Apart           week20      100
## 7 Backstreet Boys, The The One             week15      100
## 8 Carter, Aaron       Aaron's Party (Come Get It) week15      100
## 9 No Doubt            Simple Kind Of Life   week12      100
## 10 Lawrence, Tracy     Lessons Learned       week20      100
## # ... with 5,296 more rows
```

Unless we explicitly set `.by_group = TRUE`:

```
billboard_1 %>% group_by(track) %>% [c("artist", "track", "week",
"position")] %>% arrange(desc(position), .by_group = TRUE)
```

```
## # A tibble: 5,306 x 4
## # Groups:   track [316]
##   artist track           week position
##   <chr> <chr>           <chr>    <int>
## 1 Nelly (Hot S**t) Country Grammar week1      100
## 2 Nelly (Hot S**t) Country Grammar week2       99
## 3 Nelly (Hot S**t) Country Grammar week3       96
## 4 Nelly (Hot S**t) Country Grammar week4       76
## 5 Nelly (Hot S**t) Country Grammar week5       55
## 6 Nelly (Hot S**t) Country Grammar week34      49
## 7 Nelly (Hot S**t) Country Grammar week6       37
## 8 Nelly (Hot S**t) Country Grammar week11      37
## 9 Nelly (Hot S**t) Country Grammar week32      37
## 10 Nelly (Hot S**t) Country Grammar week10      36
## # ... with 5,296 more rows
```

*** Grouped `mutate()` and `filter()`***

Most useful in conjunction with aggregation and window functions:

```
billboard_1 %>% group_by(track) %>% [c("track", "week", "position")] %>%
mutate("highest position" = cummin(position)) %>% ungroup %>% filter(track
%in% c("This Time Around", "American Pie")) %>% head(n = 20)
```

```
## # A tibble: 16 x 4
##   track           week position `highest position`
##   <chr>           <chr>    <int>          <int>
## 1 This Time Around week1      22            22
## 2 This Time Around week2      22            22
## 3 This Time Around week3      20            20
## 4 This Time Around week4      45            20
## 5 This Time Around week5      87            20
## 6 This Time Around week6      71            20
## 7 This Time Around week7      95            20
## 8 American Pie     week1      43            43
## 9 American Pie     week2      35            35
## 10 American Pie     week3      29            29
## 11 American Pie     week4      29            29
## 12 American Pie     week5      33            29
```



```
## 13 American Pie      week6      32      29
## 14 American Pie      week7      40      29
## 15 American Pie      week8      58      29
## 16 American Pie      week9      88      29

billboard_1 %>% group_by(track) %>% .[c("track", "week", "position")] %>%
  filter(position < lag(position)) %>% ungroup %>% filter(track %in% c("This
  Time Around", "American Pie")) %>% head(n = 20)

## # A tibble: 5 x 3
##   track      week position
##   <chr>      <chr>    <int>
## 1 This Time Around week3      20
## 2 This Time Around week6      71
## 3 American Pie     week2      35
## 4 American Pie     week3      29
## 5 American Pie     week6      32
```

[Task 6: Practicing the dplyr verbs with flights data]

Read in the data from `flights.csv` and `airports.csv` to create two tibbles named `flights` and `airports`:

```
flights <- read_csv("flights.csv", col_names = TRUE, col_types = cols())
airports <- read_csv("airports.csv", col_names = TRUE, col_types = cols())
glimpse(flights) # glimpse() is a transposed version of print

## Rows: 336,776
## Columns: 19
## $ year      <dbl> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013,
##             2013, ...
## $ month     <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
##             1, ...
## $ day       <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
##             1, ...
## $ dep_time  <dbl> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558,
##             558,...
## $ sched_dep_time <dbl> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600,
##             600,...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -
##             2, -...
## $ arr_time  <dbl> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753,
##             849...
## $ sched_arr_time <dbl> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745,
##             851...
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3,
##             7, -...
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV",
##             "B6", ...
```

```
## $ flight      <dbl> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79,
301, ...
## $ tailnum     <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN",
"N39..."
## $ origin      <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR",
"LGA"...
## $ dest        <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL",
"IAD"...
## $ air_time    <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138,
149, ...
## $ distance    <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944,
733,...
## $ hour        <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6,
6, ...
## $ minute      <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 0,
59, ...
## $ time_hour   <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-
01 ...
```

summary(flights)

```
##      year      month      day      dep_time
sched_dep_time
## Min.   :2013   Min.    : 1.000   Min.    : 1.00   Min.     :    1   Min.    :
106
## 1st Qu.:2013   1st Qu.: 4.000   1st Qu.: 8.00   1st Qu.: 907   1st Qu.:
906
## Median :2013   Median : 7.000   Median :16.00   Median :1401   Median
:1359
## Mean   :2013   Mean    : 6.549   Mean    :15.71   Mean    :1349   Mean
:1344
## 3rd Qu.:2013   3rd Qu.:10.000   3rd Qu.:23.00   3rd Qu.:1744   3rd
Qu.:1729
## Max.   :2013   Max.    :12.000   Max.    :31.00   Max.    :2400   Max.
:2359
##
##                                     NA's    :8255
##      dep_delay      arr_time      sched_arr_time      arr_delay
## Min.   : -43.00   Min.     :    1   Min.     :    1   Min.     : -86.000
## 1st Qu.:  -5.00   1st Qu.:1104   1st Qu.:1124   1st Qu.: -17.000
## Median :  -2.00   Median :1535   Median :1556   Median :  -5.000
## Mean    : 12.64   Mean    :1502   Mean    :1536   Mean     :   6.895
## 3rd Qu.: 11.00   3rd Qu.:1940   3rd Qu.:1945   3rd Qu.: 14.000
## Max.    :1301.00   Max.    :2400   Max.    :2359   Max.    :1272.000
## NA's    :8255     NA's    :8713     NA's    :9430
##      carrier      flight      tailnum      origin
## Length:336776     Min.     :    1   Length:336776     Length:336776
## Class :character   1st Qu.: 553   Class :character   Class :character
## Mode  :character   Median :1496   Mode  :character   Mode  :character
##                                     Mean    :1972
##                                     3rd Qu.:3465
```

```
##           Max.      :8500
##
##      dest          air_time      distance      hour
## Length:336776      Min.       : 20.0      Min.       : 17      Min.       : 1.00
## Class :character    1st Qu.: 82.0      1st Qu.: 502      1st Qu.: 9.00
## Mode  :character    Median :129.0      Median : 872      Median :13.00
##                               Mean  :150.7      Mean  :1040      Mean  :13.18
##                               3rd Qu.:192.0      3rd Qu.:1389      3rd Qu.:17.00
##                               Max.   :695.0      Max.   :4983      Max.   :23.00
##                               NA's   :9430
##      minute      time_hour
## Min.       : 0.00      Min.       :2013-01-01 05:00:00
## 1st Qu.: 8.00      1st Qu.:2013-04-04 13:00:00
## Median :29.00      Median :2013-07-03 10:00:00
## Mean  :26.23      Mean  :2013-07-03 05:02:36
## 3rd Qu.:44.00      3rd Qu.:2013-10-01 07:00:00
## Max.   :59.00      Max.   :2013-12-31 23:00:00
##
```

The flights dataset is about all the flights that departed from New York City (i.e. airports JFK, LGA or EWR) in 2013. In particular, the interest lies in the following variables:

- hour, minute: the hour and minute of the departure
- arr_delay: the arrival delay of the incoming plane (in minutes)
- dest: the destination

Note that several variables have **missing values**.

(a) Creating new variables

Create a new variable which encodes a given hour and minute as one decimal number, i.e. time in hours; that is, for example, 2 hours 45 minutes should be coded to 2.75 hours, because 45 minutes is $45 \text{ minutes} * (1 \text{ hour} / 60 \text{ minutes}) = 0.75 \text{ hours}$. Name this new variable `time` in `flights`.

(b) Plotting the destinations.

(b-1) Calculate the average value of the arrival delay (`arr_delay`) and the number of departing flights (`n`) for each destination (`dest`) and name the resulting data frame `delay.per.dest`.

Tips: group observations by `dest`; use the summary function `n()` to find the count for each group; when calculating the mean, you should be aware of the existence of the missing values.

(b-2) The airports dataset contains the coordinates (`lon`, `lat`) of the 1,458 airports. `faa` stands for FAA airport code.

```
airports
```

```
## # A tibble: 1,458 x 8
##   faa   name                lat    lon   alt    tz dst  tzone
##   <chr> <chr>                <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 04G   Lansdowne Airport      41.1  -80.6  1044   -5  A
America/New_Yo...
## 2 06A   Moton Field Municipal A... 32.5  -85.7   264   -6  A
America/Chicago
## 3 06C   Schaumburg Regional      42.0  -88.1   801   -6  A
America/Chicago
## 4 06N   Randall Airport         41.4  -74.4   523   -5  A
America/New_Yo...
## 5 09J   Jekyll Island Airport    31.1  -81.4    11   -5  A
America/New_Yo...
## 6 0A9   Elizabethton Municipal ... 36.4  -82.2  1593   -5  A
America/New_Yo...
## 7 0G6   Williams County Airport  41.5  -84.5   730   -5  A
America/New_Yo...
## 8 0G7   Finger Lakes Regional A... 42.9  -76.8   492   -5  A
America/New_Yo...
## 9 0P2   Shoestring Aviation Air... 39.8  -76.6  1000   -5  U
America/New_Yo...
## 10 0S9   Jefferson County Intl     48.1 -123.    108   -8  A
America/Los_An...
## # ... with 1,448 more rows
```

Merge the tibble `delay.per.dest` and `airports` in order to add the coordinates (`lon`, `lat`) of the airports to `delay.per.dest` using `left_join()` in `dplyr` (type `?left_join` to see how it works).

(b-3) Once you have `delay.per.dest` ready, run the plotting function below to create a scatter plot of the latitude against the longitude and scale the points according to the number of departing planes.

```
plot(lat ~ lon, data = delay.per.dest, pch = 19, cex = n / 6000)
```

Explain what values are represented by the size of the bubbles.

[End of Task 6]